

```
from google.colab import drive
drive.mount('/content/drive')
```

```
Mounted at /content/drive
```

```
!pip install scikit-learn-extra
```

```
Collecting scikit-learn-extra
  Downloading scikit_learn_extra-0.3.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (2.0 MB)
    2.0/2.0 MB 12.3 MB/s eta 0:00:00
Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib/python3.10/dist-packages (from scikit-learn-extra) (1.25.2)
Requirement already satisfied: scipy>=0.19.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn-extra) (1.11.4)
Requirement already satisfied: scikit-learn>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn-extra) (1.2.2)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.23.0->scikit-learn-extra) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.23.0->scikit-learn-extra) (3.2.0)
Installing collected packages: scikit-learn-extra
Successfully installed scikit-learn-extra-0.3.0
```

```
import os
```

```
# Define the path to your dataset folder
dataset_path = '/content/drive/MyDrive/Seeds'
```

```
# Check the contents of the folder
os.chdir(dataset_path)
```

```
from skimage.feature import canny
from skimage.feature import daisy
from sklearn_extra.cluster import KMedoids
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import os
import cv2
import numpy as np
from skimage import exposure
from skimage.filters import gabor
from skimage.feature import local_binary_pattern
from skimage.feature import greycomatrix, greycoprops
from skimage.feature import hog
from sklearn.preprocessing import MinMaxScaler
```

```
# Function to extract color histograms
def extract_color_histogram(image):
    hsv_image = cv2.cvtColor(image, cv2.COLOR_BGR2HSV)
    hist_hue = cv2.calcHist([hsv_image], [0], None, [256], [0, 256])
    hist_saturation = cv2.calcHist([hsv_image], [1], None, [256], [0, 256])
    hist_value = cv2.calcHist([hsv_image], [2], None, [256], [0, 256])
    return hist_hue.flatten(), hist_saturation.flatten(), hist_value.flatten()
```

```
# Function to extract Gabor filter features
def extract_gabor_features(image):
    gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    gabor_result, _ = gabor(gray_image, frequency=0.6, theta=1.5)
    return gabor_result.flatten()
```

```
# Function to extract GLCM (Gray Level Co-occurrence Matrix) features
def extract_glcm_features(image):
    gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    distances = [1, 2, 3]
    angles = [0, np.pi/4, np.pi/2, 3*np.pi/4]
    glcm = greycomatrix(gray_image, distances=distances, angles=angles, symmetric=True, normed=True)
    contrast = greycoprops(glcm, 'contrast')
    dissimilarity = greycoprops(glcm, 'dissimilarity')
    homogeneity = greycoprops(glcm, 'homogeneity')
    energy = greycoprops(glcm, 'energy')
    correlation = greycoprops(glcm, 'correlation')
    glcm_features = np.hstack((contrast.flatten(), dissimilarity.flatten(), homogeneity.flatten(), energy.flatten(), correlation.flatten()))
    return glcm_features
```

```
# Function to extract HOG (Histogram of Oriented Gradients) features using skimage
def extract_skimage_hog_features(image):
    gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    features, _ = hog(gray_image, orientations=9, pixels_per_cell=(8, 8), cells_per_block=(2, 2), visualize=True)
    return features.flatten()

# Function to extract Local Binary Pattern (LBP) features using skimage
def extract_skimage_lbp_features(image):
    gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    radius = 3
    n_points = 8 * radius
    lbp = local_binary_pattern(gray_image, n_points, radius, method='uniform')
    hist, _ = np.histogram(lbp.ravel(), bins=np.arange(0, n_points + 3), range=(0, n_points + 2))
    return hist.flatten()

# Function to extract Canny edge features
def extract_canny_features(image):
    gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    edges = canny(gray_image)
    return edges.flatten()

# Function to extract DAISY (Discrete Anisotropic Scale Space) features
# def extract_daisy_features(image):
#     gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
#     features = daisy(gray_image, step=180, radius=58, rings=2, histograms=6, orientations=8)
#     return features.flatten()

# Function to extract texture features using Haralick features
def extract_haralick_features(image):
    gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    textures = cv2.imgproc.textureFast(gray_image)
    # Extract Haralick texture features
    haralick = cv2.imgproc.haralick(textures)
    features = np.ravel(haralick)
    return features

# Function to perform advanced image processing
def advanced_image_processing(image):
    gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    equalized_image = exposure.equalize_hist(gray_image)
    clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8, 8))
    clahe_image = clahe.apply(gray_image)
    bilateral_filtered = cv2.bilateralFilter(image, 9, 75, 75)
    return equalized_image, clahe_image, bilateral_filtered
```

```

def extract_features_and_clustering(dataset_path):
    images = []
    min_height = 100
    min_width = 100

    for subdir, dirs, files in os.walk(dataset_path):
        for file in files:
            file_path = os.path.join(subdir, file)
            img = cv2.imread(file_path)
            if img.shape[0] >= min_height and img.shape[1] >= min_width:
                images.append(img)
            else:
                print(f"Ignoring {file_path} due to small dimensions.")

    color_histograms = []
    gabor_features_list = []
    glcm_features_list = []
    skimage_hog_features_list = []
    skimage_lbp_features_list = []
    canny_features_list = []
    haralick_features_list = []

    for img in images:
        hist_hue, hist_saturation, hist_value = extract_color_histogram(img)
        color_histograms.append(np.concatenate([hist_hue, hist_saturation, hist_value]))

        gabor_features = extract_gabor_features(img)
        gabor_features_list.append(gabor_features)

        glcm_features = extract_glcm_features(img)
        glcm_features_list.append(glcm_features)

        skimage_hog_features = extract_skimage_hog_features(img)
        skimage_hog_features_list.append(skimage_hog_features)

        skimage_lbp_features = extract_skimage_lbp_features(img)
        skimage_lbp_features_list.append(skimage_lbp_features)

        canny_features = extract_canny_features(img)
        canny_features_list.append(canny_features)

    # Truncate or pad lists to ensure equal lengths
    # Ensure equal lengths of lists
    min_length = min(len(color_histograms), len(gabor_features_list), len(glcm_features_list),
                      len(skimage_hog_features_list), len(skimage_lbp_features_list), len(canny_features_list))

    color_histograms = [feature[:min_length] for feature in color_histograms]
    gabor_features_list = [feature[:min_length] for feature in gabor_features_list]
    glcm_features_list = [feature[:min_length] for feature in glcm_features_list]
    skimage_hog_features_list = [feature[:min_length] for feature in skimage_hog_features_list]
    skimage_lbp_features_list = [feature[:min_length] for feature in skimage_lbp_features_list]
    canny_features_list = [feature[:min_length] for feature in canny_features_list]

    # Concatenate the features
    features = np.concatenate([color_histograms, gabor_features_list, glcm_features_list,
                               skimage_hog_features_list, skimage_lbp_features_list, canny_features_list], axis=1)

    scaler = MinMaxScaler()
    features_scaled = scaler.fit_transform(features)

    k = 5
    kmedoids = KMedoids(n_clusters=k)
    clusters = kmedoids.fit_predict(features_scaled)

    return features_scaled, clusters

```

```

def encode_features(features):
    from sklearn.decomposition import PCA

    # Determine the maximum value for n_components based on the minimum of samples or features
    n_components = min(features.shape[0], features.shape[1])

    # Reduce n_components to avoid exceeding the minimum of samples or features
    n_components = min(n_components, 90) # Set a maximum value of 90 for n_components

    pca = PCA(n_components=n_components)
    encoded_features = pca.fit_transform(features)
    return encoded_features

# # Function to perform Decision Tree Classification
# def decision_tree_classification(encoded_features, clusters):
#     X_train, X_test, y_train, y_test = train_test_split(encoded_features, clusters, test_size=0.2, random_state=42)
#     tree_classifier = DecisionTreeClassifier(random_state=42)
#     tree_classifier.fit(X_train, y_train)
#     y_pred = tree_classifier.predict(X_test)
#     accuracy = accuracy_score(y_test, y_pred)
#     return accuracy

import matplotlib.pyplot as plt

from sklearn.metrics import roc_curve, confusion_matrix, roc_auc_score

# Function to plot ROC curve
def plot_roc_curve(y_test, y_pred_proba):
    fpr, tpr, _ = roc_curve(y_test, y_pred_proba[:, 1])
    plt.figure(figsize=(8, 6))
    plt.plot(fpr, tpr, label='ROC Curve')
    plt.plot([0, 1], [0, 1], 'k--', label='Random')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC Curve for Decision Tree Classifier')
    plt.legend()
    plt.show()

# Function to display confusion matrix
def display_confusion_matrix(y_test, y_pred):
    cm = confusion_matrix(y_test, y_pred)
    print("Confusion Matrix:")
    print(cm)

# Function to perform Decision Tree Classification and get metrics
def decision_tree_classification_metrics(encoded_features, clusters):
    X_train, X_test, y_train, y_test = train_test_split(encoded_features, clusters, test_size=0.2, random_state=42)
    tree_classifier = DecisionTreeClassifier(random_state=42)
    tree_classifier.fit(X_train, y_train)
    y_pred = tree_classifier.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)

    # Calculate ROC AUC score
    y_pred_proba = tree_classifier.predict_proba(X_test)
    roc_auc = roc_auc_score(y_test, y_pred_proba[:, 1])

    # Plot ROC curve
    plot_roc_curve(y_test, y_pred_proba)

    # Display confusion matrix
    display_confusion_matrix(y_test, y_pred)

    return accuracy, roc_auc

# Define the path to your dataset folder
dataset_path = '/content/drive/MyDrive/Seeds'

# Extract features and perform clustering
features_scaled, clusters = extract_features_and_clustering(dataset_path)

```

```

Ignoring /content/drive/MyDrive/Seeds/Litchi/Litchi (59).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Orange/Orange (82).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Sapodilla/Sapodilla (45).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Hog_Plum/Hog_Plum (63).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Hog_Plum/Hog_Plum (66).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Hog_Plum/Hog_Plum (64).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Hog_Plum/Hog_Plum (65).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Olive/Olive (76).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Olive/Olive (75).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Amla/Amla (24).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Amla/Amla (27).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Amla/Amla (23).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Amla/Amla (26).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Amla/Amla (25).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Black_Plum/Black_Plum (1).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Black_Plum/Black_Plum (2).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Guava/Guava (85).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Guava/Guava (84).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Guava/Guava (86).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Guava/Guava (88).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Guava/Guava (90).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Guava/Guava (89).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Guava/Guava (83).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Guava/Guava (81).jpg due to small dimensions.
Ignoring /content/drive/MyDrive/Seeds/Guava/Guava (87).jpg due to small dimensions.
Streaming output truncated to the last 5000 lines.
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:42: skimage_deprecation: Function ``greycoprops`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:42: skimage_deprecation: Function ``greycoprops`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:42: skimage_deprecation: Function ``greycoprops`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:42: skimage_deprecation: Function ``greycoprops`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:35: skimage_deprecation: Function ``greycomatrix`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:42: skimage_deprecation: Function ``greycoprops`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:42: skimage_deprecation: Function ``greycoprops`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:42: skimage_deprecation: Function ``greycoprops`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:42: skimage_deprecation: Function ``greycoprops`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:42: skimage_deprecation: Function ``greycoprops`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:35: skimage_deprecation: Function ``greycomatrix`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:42: skimage_deprecation: Function ``greycoprops`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:42: skimage_deprecation: Function ``greycoprops`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:42: skimage_deprecation: Function ``greycoprops`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:42: skimage_deprecation: Function ``greycoprops`` is deprec
removed_version='1.0')
/usr/local/lib/python3.10/dist-packages/skimage/feature/__init__.py:42: skimage_deprecation: Function ``greycoprops`` is deprec
removed_version='1.0')

```

```

# Encode features
encoded_features = encode_features(features_scaled)

```

```

# # Perform Decision Tree Classification
# accuracy = decision_tree_classification(encoded_features, clusters)
# print("Accuracy of Decision Tree Classifier using K-medoids clusters:", accuracy)

```

```

from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay

def decision_tree_classification_metrics(encoded_features, clusters):
    X_train, X_test, y_train, y_test = train_test_split(encoded_features, clusters, test_size=0.2, random_state=42)
    tree_classifier = DecisionTreeClassifier(random_state=42)
    tree_classifier.fit(X_train, y_train)
    y_pred = tree_classifier.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)


    # Calculate confusion matrix
    cm = confusion_matrix(y_test, y_pred)

    # Display confusion matrix
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=np.unique(y_test))
    disp.plot(cmap='Blues')
    plt.title('Confusion Matrix for Decision Tree Classifier')
    plt.show()

    return accuracy

# Perform Decision Tree Classification and plot confusion matrix
dec_accuracy = decision_tree_classification_metrics(encoded_features, clusters)
print("Accuracy of Decision Tree Classifier using K-medoids clusters:", dec_accuracy)

```



Accuracy of Decision Tree Classifier using K-medoids clusters: 0.8450704225352113

```

from sklearn.svm import SVC

from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay

def svm_classification_metrics(encoded_features, clusters):
    X_train, X_test, y_train, y_test = train_test_split(encoded_features, clusters, test_size=0.2, random_state=42)
    svm_classifier = SVC(kernel='linear', probability=True, random_state=42)
    svm_classifier.fit(X_train, y_train)
    y_pred = svm_classifier.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)


    # Calculate confusion matrix
    cm = confusion_matrix(y_test, y_pred)

    # Display confusion matrix
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=np.unique(y_test))
    disp.plot(cmap='Blues')
    plt.title('Confusion Matrix for SVM Classifier')
    plt.show()

    return accuracy

# Perform SVM Classification and plot confusion matrix
svm_accuracy = svm_classification_metrics(encoded_features, clusters)
print("Accuracy of SVM Classifier using K-medoids clusters:", svm_accuracy)

```



Accuracy of SVM Classifier using K-medoids clusters: 0.8985915492957747

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier


# Function to perform Random Forest Classification
def random_forest_classification_metrics(encoded_features, clusters):
    X_train, X_test, y_train, y_test = train_test_split(encoded_features, clusters, test_size=0.2, random_state=42)
    rf_classifier = RandomForestClassifier(random_state=42)
    rf_classifier.fit(X_train, y_train)
    y_pred = rf_classifier.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)

    # Calculate confusion matrix
    cm = confusion_matrix(y_test, y_pred)

    # Display confusion matrix
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=np.unique(y_test))
    disp.plot(cmap='Blues')
    plt.title('Confusion Matrix for Random Forest Classifier')
    plt.show()

    return accuracy

# Perform Random Forest Classification and plot confusion matrix
rf_accuracy = random_forest_classification_metrics(encoded_features, clusters)
print("Accuracy of Random Forest Classifier using K-medoids clusters:", rf_accuracy)
```



Accuracy of Random Forest Classifier using K-medoids clusters: 0.8816901408450705

```
# Function to perform k-Nearest Neighbors (KNN) Classification
def knn_classification_metrics(encoded_features, clusters):
    X_train, X_test, y_train, y_test = train_test_split(encoded_features, clusters, test_size=0.2, random_state=42)
    knn_classifier = KNeighborsClassifier()
    knn_classifier.fit(X_train, y_train)
    y_pred = knn_classifier.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)


    # Calculate confusion matrix
    cm = confusion_matrix(y_test, y_pred)

    # Display confusion matrix
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=np.unique(y_test))
    disp.plot(cmap='Blues')
    plt.title('Confusion Matrix for k-Nearest Neighbors (KNN) Classifier')
    plt.show()

    return accuracy

# Perform k-Nearest Neighbors (KNN) Classification and plot confusion matrix
knn_accuracy = knn_classification_metrics(encoded_features, clusters)
print("Accuracy of k-Nearest Neighbors (KNN) Classifier using K-medoids clusters:", knn_accuracy)
```

Confusion Matrix for k-Nearest Neighbors (KNN) Classifier



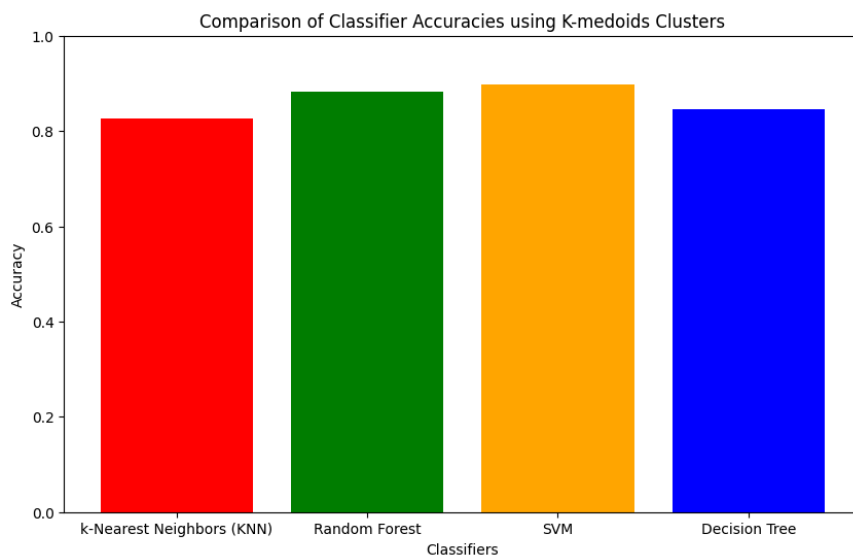
```
import matplotlib.pyplot as plt
```

```
# Accuracy values
accuracies = [knn_accuracy, rf_accuracy, svm_accuracy, dec_accuracy]

# Classifier names
classifiers = ['k-Nearest Neighbors (KNN)', 'Random Forest', 'SVM', 'Decision Tree']

# Plotting the bar graph
plt.figure(figsize=(10, 6))
plt.bar(classifiers, accuracies, color=['red', 'green', 'orange', 'blue'])
plt.xlabel('Classifiers')
plt.ylabel('Accuracy')
plt.title('Comparison of Classifier Accuracies using K-medoids Clusters')
plt.ylim(0, 1) # Setting y-axis limit to better visualize differences
plt.show()
```





```
import numpy as np
from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score, precision_recall_curve, roc_curve

def calculate_metrics(y_test, y_pred, y_pred_proba):
    # Precision
    precision = precision_score(y_test, y_pred, average='weighted')

    # Recall
    recall = recall_score(y_test, y_pred, average='weighted')

    # F1 Score
    f1 = f1_score(y_test, y_pred, average='weighted')

    return precision, recall, f1

def decision_tree_classification_metrics(encoded_features, clusters):
    X_train, X_test, y_train, y_test = train_test_split(encoded_features, clusters, test_size=0.2, random_state=42)
    tree_classifier = DecisionTreeClassifier(random_state=42)
    tree_classifier.fit(X_train, y_train)
    y_pred = tree_classifier.predict(X_test)

    # Calculate probabilities for positive class
    y_pred_proba = tree_classifier.predict_proba(X_test)

    # Calculate additional metrics
    precision, recall, f1 = calculate_metrics(y_test, y_pred, y_pred_proba)

    print("Precision:", precision)
    print("Recall:", recall)
    print("F1 Score:", f1)

decision_tree_classification_metrics(encoded_features, clusters)

Precision: 0.8419296118996915
Recall: 0.8450704225352113
F1 Score: 0.8432251946639185
```

```
def svm_classification_metrics(encoded_features, clusters):
    X_train, X_test, y_train, y_test = train_test_split(encoded_features, clusters, test_size=0.2, random_state=42)
    svm_classifier = SVC(kernel='linear', probability=True, random_state=42)
    svm_classifier.fit(X_train, y_train)
    y_pred = svm_classifier.predict(X_test)

    # Calculate probabilities for positive class
    y_pred_proba = svm_classifier.predict_proba(X_test)

    # Calculate additional metrics
    precision, recall, f1 = calculate_metrics(y_test, y_pred, y_pred_proba)

    print("Precision:", precision)
    print("Recall:", recall)
    print("F1 Score:", f1)

# Perform SVM Classification and compute all metrics
svm_classification_metrics(encoded_features, clusters)
```

```
Precision: 0.8968943797342201
Recall: 0.8985915492957747
F1 Score: 0.8942339272987956
```

```
def random_forest_classification_metrics(encoded_features, clusters):
    X_train, X_test, y_train, y_test = train_test_split(encoded_features, clusters, test_size=0.2, random_state=42)
    rf_classifier = RandomForestClassifier(random_state=42)
    rf_classifier.fit(X_train, y_train)
    y_pred = rf_classifier.predict(X_test)

    # Calculate probabilities for positive class
    y_pred_proba = rf_classifier.predict_proba(X_test)

    # Calculate additional metrics
    precision, recall, f1 = calculate_metrics(y_test, y_pred, y_pred_proba)

    print("Precision:", precision)
    print("Recall:", recall)
    print("F1 Score:", f1)
```

```
random_forest_classification_metrics(encoded_features, clusters)
```

```
Precision: 0.8798728591168977
Recall: 0.8816901408450705
F1 Score: 0.8757002915408361
```

```
def knn_classification_metrics(encoded_features, clusters):
    X_train, X_test, y_train, y_test = train_test_split(encoded_features, clusters, test_size=0.2, random_state=42)
    knn_classifier = KNeighborsClassifier()
    knn_classifier.fit(X_train, y_train)
    y_pred = knn_classifier.predict(X_test)

    # Calculate probabilities for positive class
    y_pred_proba = knn_classifier.predict_proba(X_test)

    # Calculate additional metrics
    precision, recall, f1 = calculate_metrics(y_test, y_pred, y_pred_proba)

    print("Precision:", precision)
    print("Recall:", recall)
    print("F1 Score:", f1)
```

```
knn_classification_metrics(encoded_features, clusters)
```

```
Precision: 0.8195876788488435
Recall: 0.8253521126760563
F1 Score: 0.8210066047072261
```

```
import matplotlib.pyplot as plt


# Accuracy values
accuracies = [knn_accuracy, rf_accuracy, svm_accuracy, dec_accuracy]

# Classifier names
classifiers = ['k-Nearest Neighbors (KNN)', 'Random Forest', 'SVM', 'Decision Tree']

# Plotting the bar graph
plt.figure(figsize=(10, 6))
bars = plt.bar(classifiers, accuracies, color=['red', 'green', 'orange', 'blue'])
plt.xlabel('Classifiers')
plt.ylabel('Accuracy')
plt.title('Comparison of Classifier Accuracies using K-medoids Clusters')
plt.ylim(0, 1) # Setting y-axis limit to better visualize differences

# Adding text annotations on top of each bar
for bar, accuracy in zip(bars, accuracies):
    plt.text(bar.get_x() + bar.get_width() / 2, bar.get_height() + 0.01, f'{accuracy:.2%}', ha='center', va='bottom')

plt.show()
```



```

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report

# Split the data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(encoded_features, clusters, test_size=0.2, random_state=42)

# Initialize Random Forest classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)

# Train the classifier
rf_classifier.fit(X_train, y_train)

# Predict on the test set
y_pred = rf_classifier.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Generate classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Display confusion matrix
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(cm)

```

Accuracy: 0.8816901408450705

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.99	0.97	191
1	0.88	0.71	0.79	31
2	0.87	0.54	0.67	24
3	0.78	0.90	0.84	88
4	0.62	0.48	0.54	21
accuracy			0.88	355
macro avg	0.82	0.72	0.76	355
weighted avg	0.88	0.88	0.88	355

Confusion Matrix:

```

[[189  0  0  2  0]
 [ 0 22  2  2  5]
 [ 0  1 13  9  1]
 [ 9  0  0 79  0]
 [ 0  2  0  9 10]]

```

```

from sklearn.ensemble import VotingClassifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(encoded_features, clusters, test_size=0.2, random_state=42)

# Define individual classifiers
svm_classifier = SVC(probability=True)
rf_classifier = RandomForestClassifier()
knn_classifier = KNeighborsClassifier()

# Define the ensemble model using voting
voting_classifier = VotingClassifier(estimators=[('svm', svm_classifier), ('rf', rf_classifier), ('knn', knn_classifier)], voting='soft')

# Train the ensemble model
voting_classifier.fit(X_train, y_train)

# Predict the labels
y_pred = voting_classifier.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)

# Calculate ROC AUC score
y_pred_proba = voting_classifier.predict_proba(X_test)

# Display confusion matrix and ROC curve
display_confusion_matrix(y_test, y_pred)

# Print accuracy and ROC AUC score
print("Accuracy:", accuracy)

```

```

Confusion Matrix:
[[189  0  0  2  0]
 [ 0 19  1  5  6]
 [ 0  5  9  9  1]
 [ 8  3  2 74  1]
 [ 0  4  1  7  9]]
Accuracy: 0.8450704225352113

```

```

from sklearn.ensemble import GradientBoostingClassifier

# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(encoded_features, clusters, test_size=0.2, random_state=42)

# Initialize Gradient Boosting classifier
gb_classifier = GradientBoostingClassifier(n_estimators=100, random_state=42)

# Train the classifier
gb_classifier.fit(X_train, y_train)

# Predict the labels for test set
y_pred = gb_classifier.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Display confusion matrix
def display_confusion_matrix(y_test, y_pred):
    cm = confusion_matrix(y_test, y_pred)
    print("Confusion Matrix:")
    print(cm)

# Display confusion matrix
display_confusion_matrix(y_test, y_pred)

```

```

Accuracy: 0.9126760563380282
Confusion Matrix:
[[189  0  0  2  0]
 [ 0 26  0  0  5]
 [ 0  1 16  7  0]
 [ 8  0  0 80  0]
 [ 0  1  0  7 13]]

```

```

from sklearn.ensemble import AdaBoostClassifier

# Initialize AdaBoost classifier
adaboost_classifier = AdaBoostClassifier(n_estimators=100, random_state=42)

# Train the classifier
adaboost_classifier.fit(X_train, y_train)

# Predict the labels for test set
y_pred = adaboost_classifier.predict(X_test)

```


```

from sklearn.metrics import accuracy_score, confusion_matrix

# Calculate accuracy
accuracy_adaboost = accuracy_score(y_test, y_pred)
print("AdaBoost Accuracy:", accuracy_adaboost)

# Compute confusion matrix
conf_matrix_adaboost = confusion_matrix(y_test, y_pred)
print("AdaBoost Confusion Matrix:")
print(conf_matrix_adaboost)

```

 AdaBoost Accuracy: 0.6  
 AdaBoost Confusion Matrix:  

```

[[125  0  0 66  0]
 [  0  3  9 19  0]
 [  0  0 17  6  1]
 [  1  1 19 66  1]
 [  0  2  8  9  2]]

```

```

from sklearn.ensemble import BaggingClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix
import numpy as np

# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(encoded_features, clusters, test_size=0.2, random_state=42)

# Initialize Bagging classifier without specifying base_estimator
bagging_classifier = BaggingClassifier(n_estimators=100, random_state=42)

# Train the classifier
bagging_classifier.fit(X_train, y_train)

# Predict the labels for test set
y_pred = bagging_classifier.predict(X_test)

# Calculate accuracy
accuracy_bagging = accuracy_score(y_test, y_pred)
print("Bagging Accuracy:", accuracy_bagging)

# Compute confusion matrix
conf_matrix_bagging = confusion_matrix(y_test, y_pred)
print("Bagging Confusion Matrix:")
print(conf_matrix_bagging)

```

Bagging Accuracy: 0.8901408450704226  
 Bagging Confusion Matrix:  

```

[[184  0  0  7  0]
 [  0 26  0  0  5]
 [  0  2 14  7  1]
 [  8  0  0 80  0]
 [  0  2  0  7 12]]

```

```

from sklearn.ensemble import ExtraTreesClassifier

# Initialize Extra Trees classifier
extra_trees_classifier = ExtraTreesClassifier(n_estimators=100, random_state=42)

# Train the classifier
extra_trees_classifier.fit(X_train, y_train)

# Predict the labels for test set
y_pred = extra_trees_classifier.predict(X_test)
# Calculate accuracy
accuracy_extra_trees = accuracy_score(y_test, y_pred)
print("Extra Trees Accuracy:", accuracy_extra_trees)

# Compute confusion matrix
conf_matrix_extra_trees = confusion_matrix(y_test, y_pred)
print("Extra Trees Confusion Matrix:")
print(conf_matrix_extra_trees)

```

```

Extra Trees Accuracy: 0.8422535211267606
Extra Trees Confusion Matrix:
[[188  0  0  3  0]
 [  0 23  0  3  5]
 [  1  2 11  8  2]
 [ 19  0  0 69  0]
 [  1  2  0 10  8]]

```

```

import lightgbm as lgb

# Convert data to LightGBM Dataset format
train_data = lgb.Dataset(X_train, label=y_train)
test_data = lgb.Dataset(X_test)

# Define parameters for LightGBM
params = {
    'objective': 'binary',
    'metric': 'auc',
    'num_leaves': 31,
    'learning_rate': 0.05,
    'feature_fraction': 0.9,
    'bagging_fraction': 0.8,
    'bagging_freq': 5,
    'verbose': -1,
    'seed': 42
}

# Train the model
lgb_classifier = lgb.train(params, train_data, num_boost_round=100)

# Predict probabilities for test set
y_pred_proba = lgb_classifier.predict(X_test)
y_pred = (y_pred_proba > 0.5).astype(int)

# Calculate accuracy
accuracy_lightgbm = accuracy_score(y_test, y_pred)
print("LightGBM Accuracy:", accuracy_lightgbm)

# Compute confusion matrix
conf_matrix_lightgbm = confusion_matrix(y_test, y_pred)
print("LightGBM Confusion Matrix:")
print(conf_matrix_lightgbm)

```

```

LightGBM Accuracy: 0.6197183098591549
LightGBM Confusion Matrix:
[[189  2  0  0  0]
 [  0 31  0  0  0]
 [  0 24  0  0  0]
 [  6 82  0  0  0]
 [  0 21  0  0  0]]

```

```
!pip install catboost
```

```

Collecting catboost
  Downloading catboost-1.2.3-cp310-cp310-manylinux2014_x86_64.whl (98.5 MB)
    ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 98.5/98.5 MB 939.7 kB/s eta 0:00:00
Requirement already satisfied: graphviz in /usr/local/lib/python3.10/dist-packages (from catboost) (0.20.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from catboost) (3.7.1)
Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.10/dist-packages (from catboost) (1.25.2)
Requirement already satisfied: pandas>=0.24 in /usr/local/lib/python3.10/dist-packages (from catboost) (1.5.3)

```

```
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from catboost) (1.11.4)
Requirement already satisfied: plotly in /usr/local/lib/python3.10/dist-packages (from catboost) (5.15.0)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from catboost) (1.16.0)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.24->catboost) (2.8)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.24->catboost) (2023.4)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (1.2.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (4.50.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (1.4.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (24.0)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (3.1.2)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from plotly->catboost) (8.2.3)
Installing collected packages: catboost
Successfully installed catboost-1.2.3
```

```
from catboost import CatBoostClassifier
```

```
# Initialize CatBoost classifier
```

```
catboost_classifier = CatBoostClassifier(iterations=100, random_seed=42)
```

```
# Train the classifier
```

```
catboost_classifier.fit(X_train, y_train)
```

```
# Predict the labels for test set
```

```
y_pred = catboost_classifier.predict(X_test)
```

```
27: learn: 0.1160075 total: 3.1s remaining: 7.98s
28: learn: 0.1131135 total: 3.23s remaining: 7.9s
29: learn: 0.1086183 total: 3.38s remaining: 7.88s
30: learn: 0.1070518 total: 3.55s remaining: 7.89s
31: learn: 0.1042950 total: 3.63s remaining: 7.72s
32: learn: 0.1006826 total: 3.72s remaining: 7.55s
33: learn: 0.0964659 total: 3.8s remaining: 7.38s
34: learn: 0.0917384 total: 3.89s remaining: 7.22s
35: learn: 0.0895130 total: 3.97s remaining: 7.06s
36: learn: 0.0872792 total: 4.06s remaining: 6.91s
37: learn: 0.0847562 total: 4.15s remaining: 6.78s
38: learn: 0.0831520 total: 4.24s remaining: 6.63s
39: learn: 0.0814695 total: 4.32s remaining: 6.48s
40: learn: 0.0777011 total: 4.41s remaining: 6.34s
41: learn: 0.0747309 total: 4.51s remaining: 6.23s
42: learn: 0.0730516 total: 4.6s remaining: 6.1s
43: learn: 0.0712693 total: 4.69s remaining: 5.97s
44: learn: 0.0686978 total: 4.77s remaining: 5.83s
45: learn: 0.0671743 total: 4.86s remaining: 5.71s
46: learn: 0.0662859 total: 4.96s remaining: 5.59s
47: learn: 0.0650100 total: 5.05s remaining: 5.47s
48: learn: 0.0631617 total: 5.14s remaining: 5.35s
49: learn: 0.0622850 total: 5.22s remaining: 5.22s
50: learn: 0.0613425 total: 5.31s remaining: 5.1s
51: learn: 0.0599545 total: 5.4s remaining: 4.98s
52: learn: 0.0574355 total: 5.49s remaining: 4.87s
53: learn: 0.0570203 total: 5.59s remaining: 4.76s
54: learn: 0.0548541 total: 5.68s remaining: 4.65s
55: learn: 0.0528318 total: 5.77s remaining: 4.54s
56: learn: 0.0511628 total: 5.86s remaining: 4.42s
57: learn: 0.0502654 total: 5.94s remaining: 4.3s
58: learn: 0.0499014 total: 6.03s remaining: 4.19s
59: learn: 0.0489920 total: 6.11s remaining: 4.08s
60: learn: 0.0481675 total: 6.25s remaining: 3.99s
61: learn: 0.0474997 total: 6.4s remaining: 3.92s
62: learn: 0.0470287 total: 6.58s remaining: 3.86s
63: learn: 0.0458852 total: 6.72s remaining: 3.78s
64: learn: 0.0448398 total: 6.89s remaining: 3.71s
65: learn: 0.0439723 total: 7.05s remaining: 3.63s
66: learn: 0.0435796 total: 7.18s remaining: 3.54s
67: learn: 0.0423085 total: 7.36s remaining: 3.46s
68: learn: 0.0414019 total: 7.56s remaining: 3.4s
69: learn: 0.0408835 total: 7.79s remaining: 3.34s
70: learn: 0.0404868 total: 8.02s remaining: 3.27s
71: learn: 0.0396479 total: 8.29s remaining: 3.22s
72: learn: 0.0387112 total: 8.54s remaining: 3.16s
73: learn: 0.0381882 total: 8.88s remaining: 3.12s
74: learn: 0.0372681 total: 9.11s remaining: 3.04s
75: learn: 0.0371015 total: 9.34s remaining: 2.95s
76: learn: 0.0365641 total: 9.64s remaining: 2.88s
77: learn: 0.0359164 total: 9.88s remaining: 2.79s
78: learn: 0.0355426 total: 10.2s remaining: 2.7s
79: learn: 0.0347783 total: 10.4s remaining: 2.61s
80: learn: 0.0339271 total: 10.6s remaining: 2.5s
81: learn: 0.0335851 total: 10.9s remaining: 2.39s
-- --
```